Semantic masking (sync-only) vs. TAP-Foley

300 videos × 10 targets (3,000 pairs); NFE matched within each backbone (MMAudio 100, Omni2Sound 400)

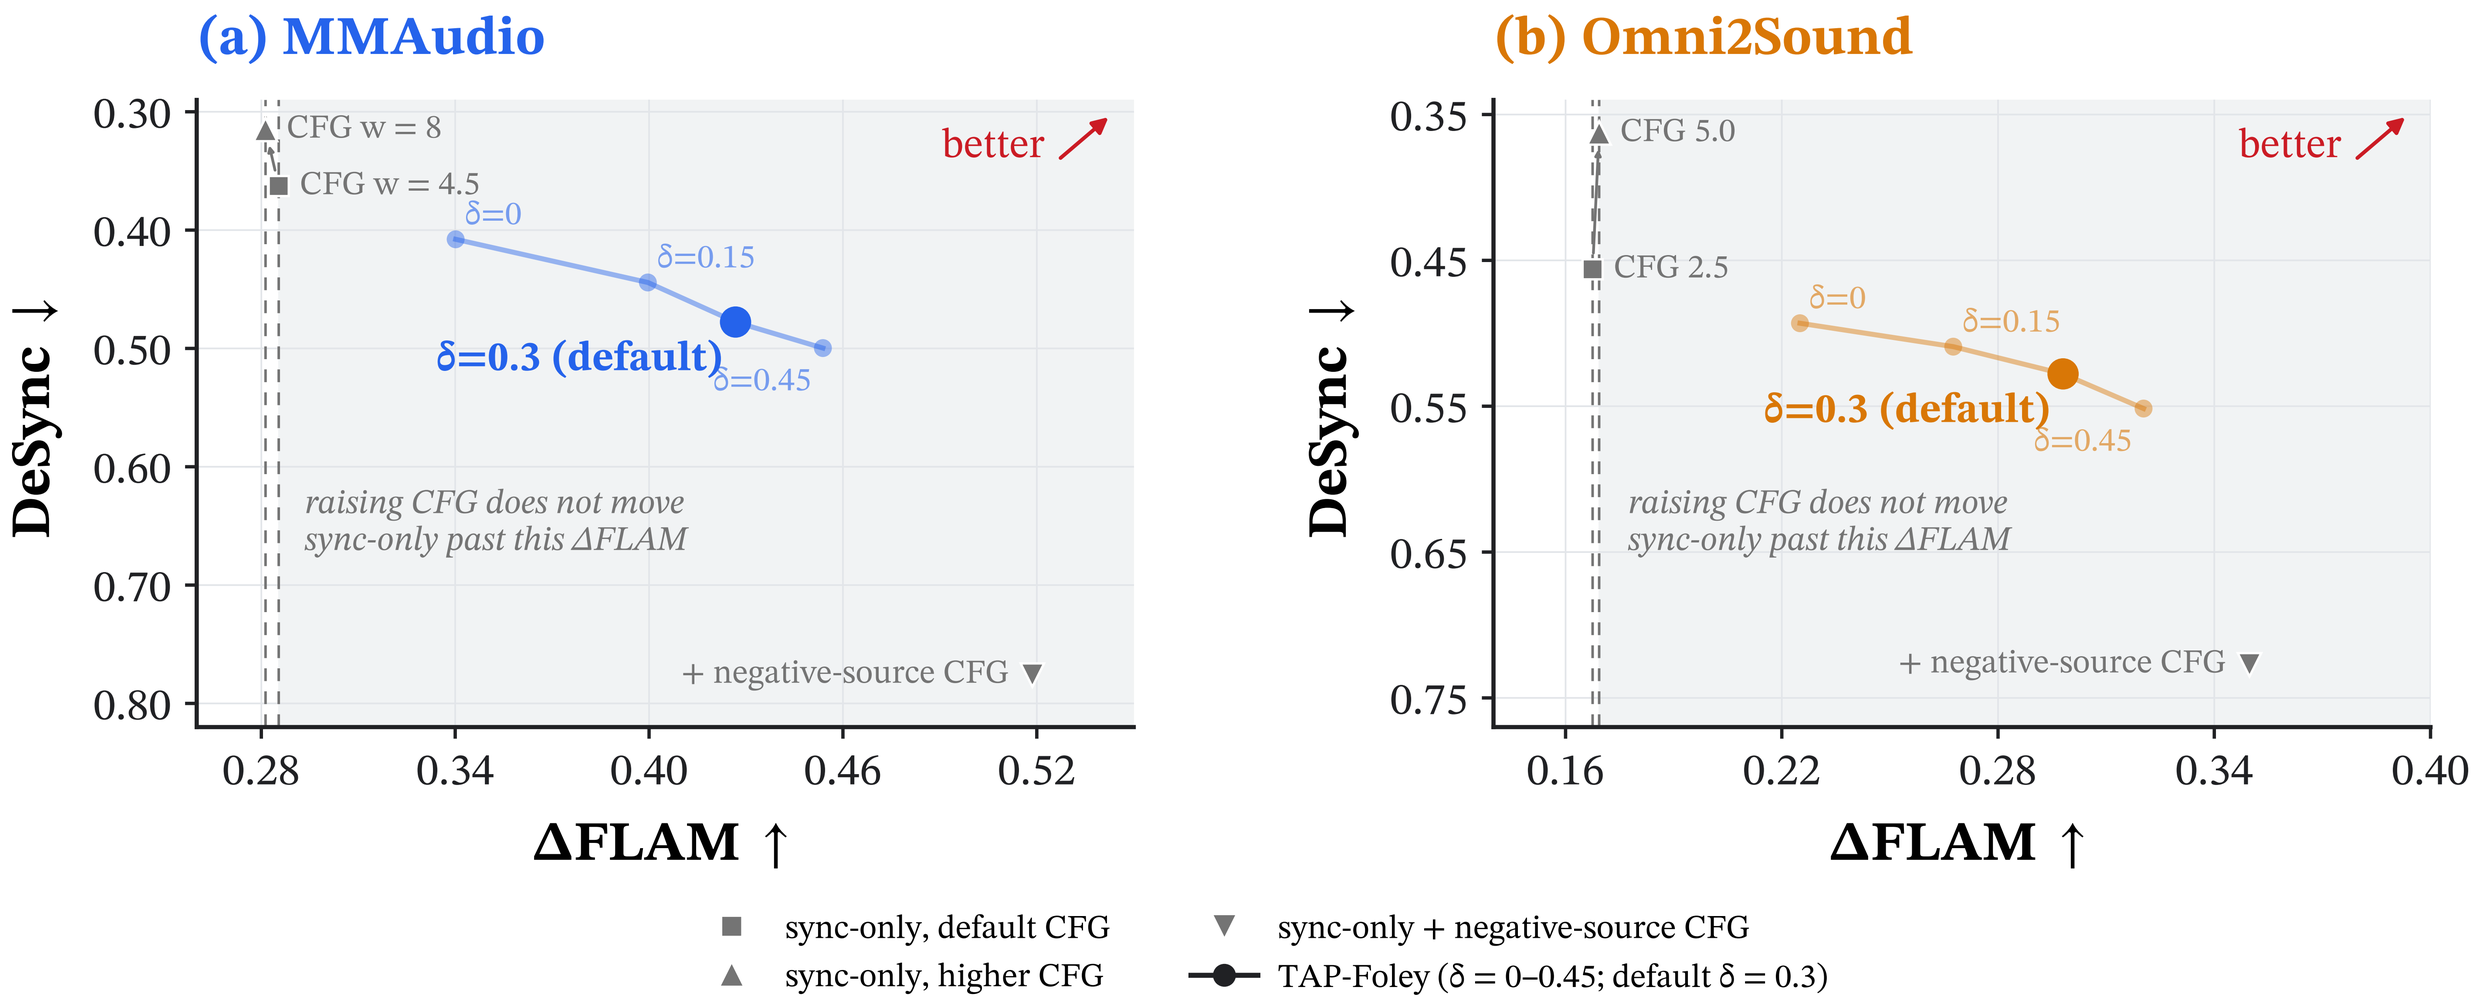

Sync-only: semantic (CLIP) video features masked, synchronization features kept. Colored marks follow the backbone color. Shaded: ΔFLAM beyond the sync-only line.