Semantic masking (sync-only) vs. TAP-Foley

held-out 300 videos × 10 targets (3,000 pairs), NFE 100 for every system

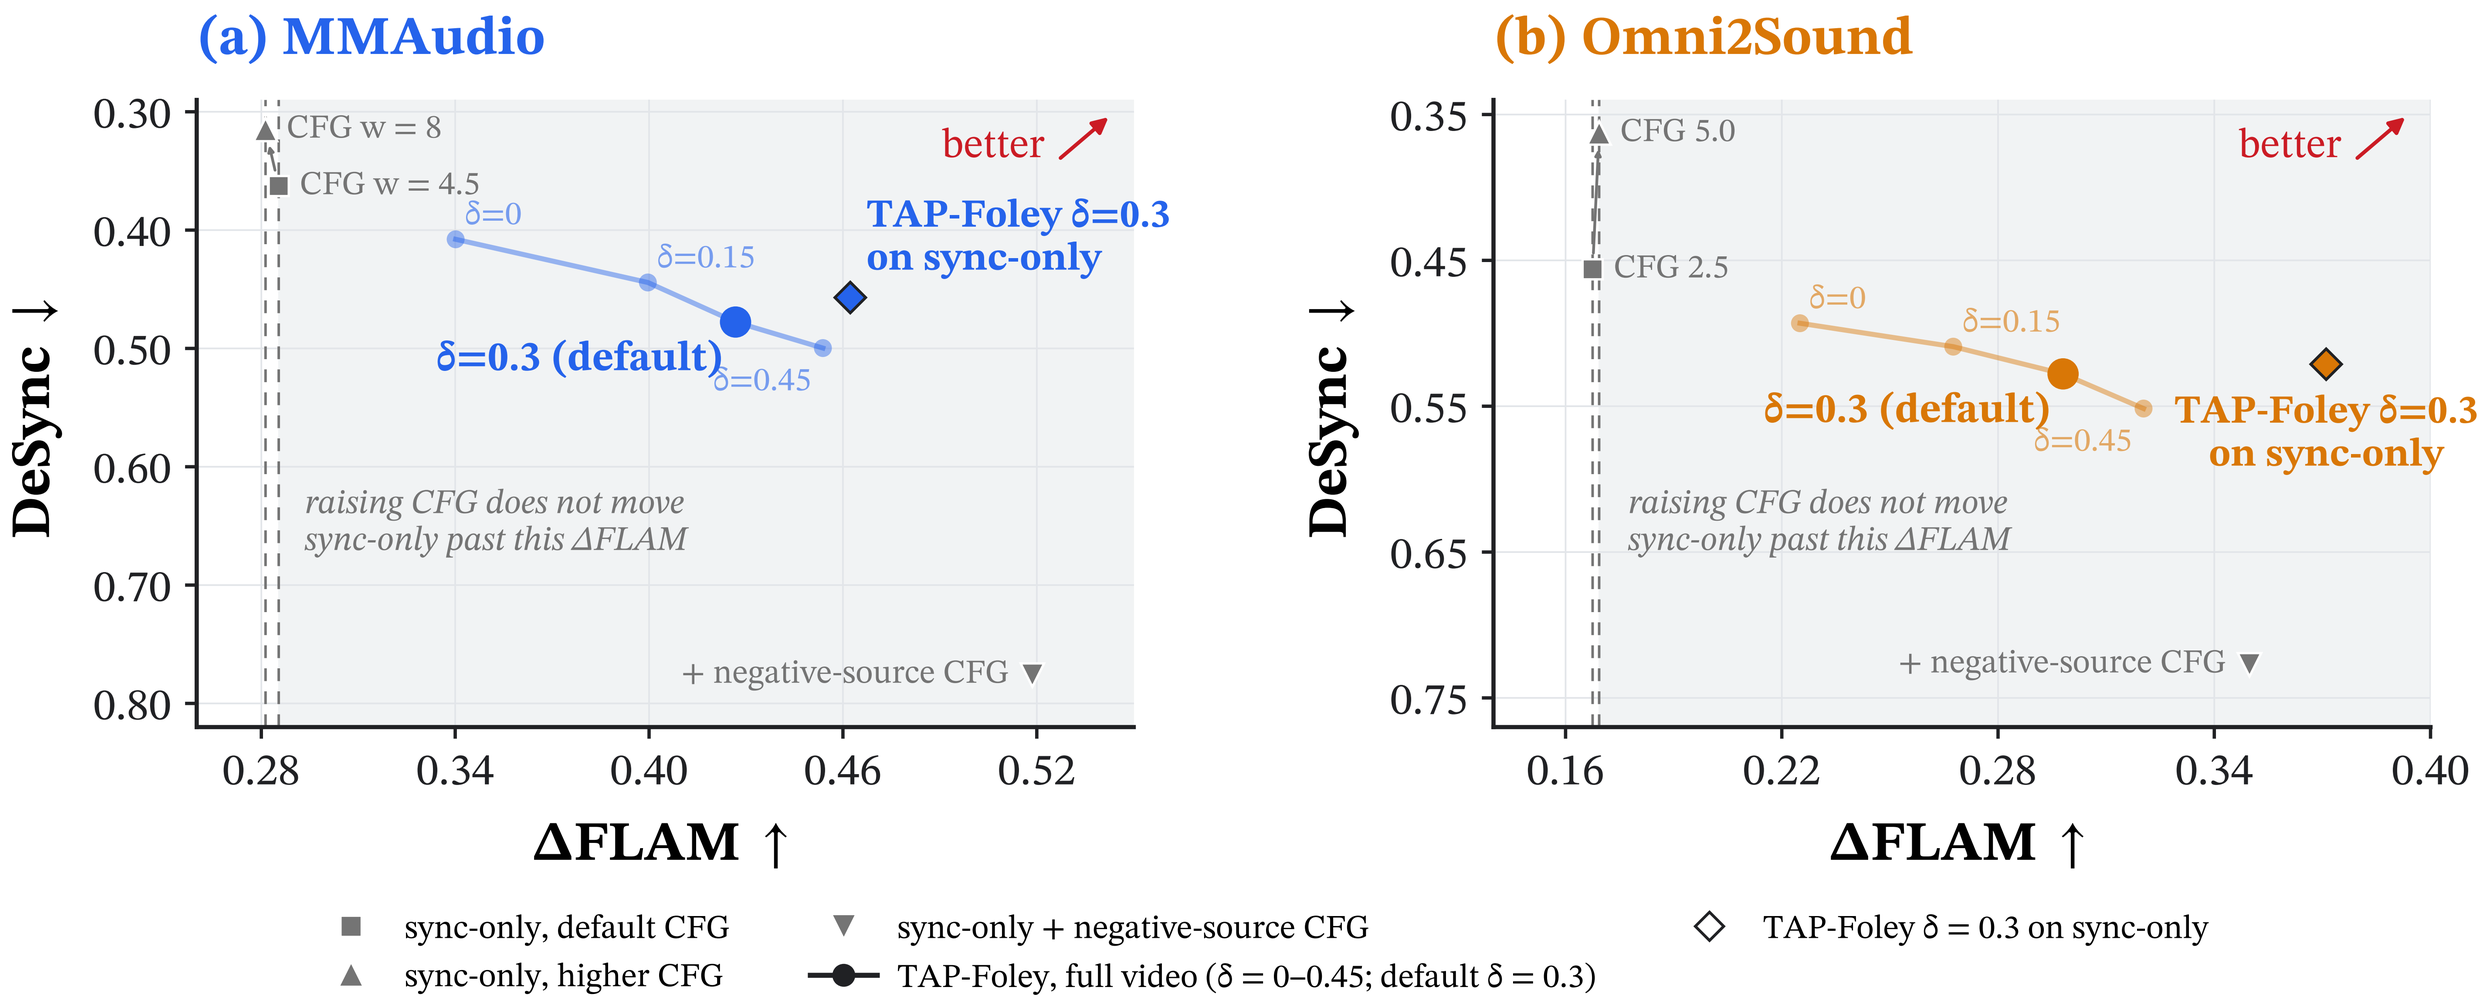

Sync-only: semantic (CLIP) video features masked, synchronization features kept. Colored marks follow the backbone color. Shaded: ΔFLAM beyond the sync-only line.