Semantic masking (sync-only) vs. TAP-Foley

Full benchmark: 1,189 videos × 10 targets (11,890 pairs); NFE matched within each backbone (MMAudio 100, Omni2Sound 400, AC-Foley 100)

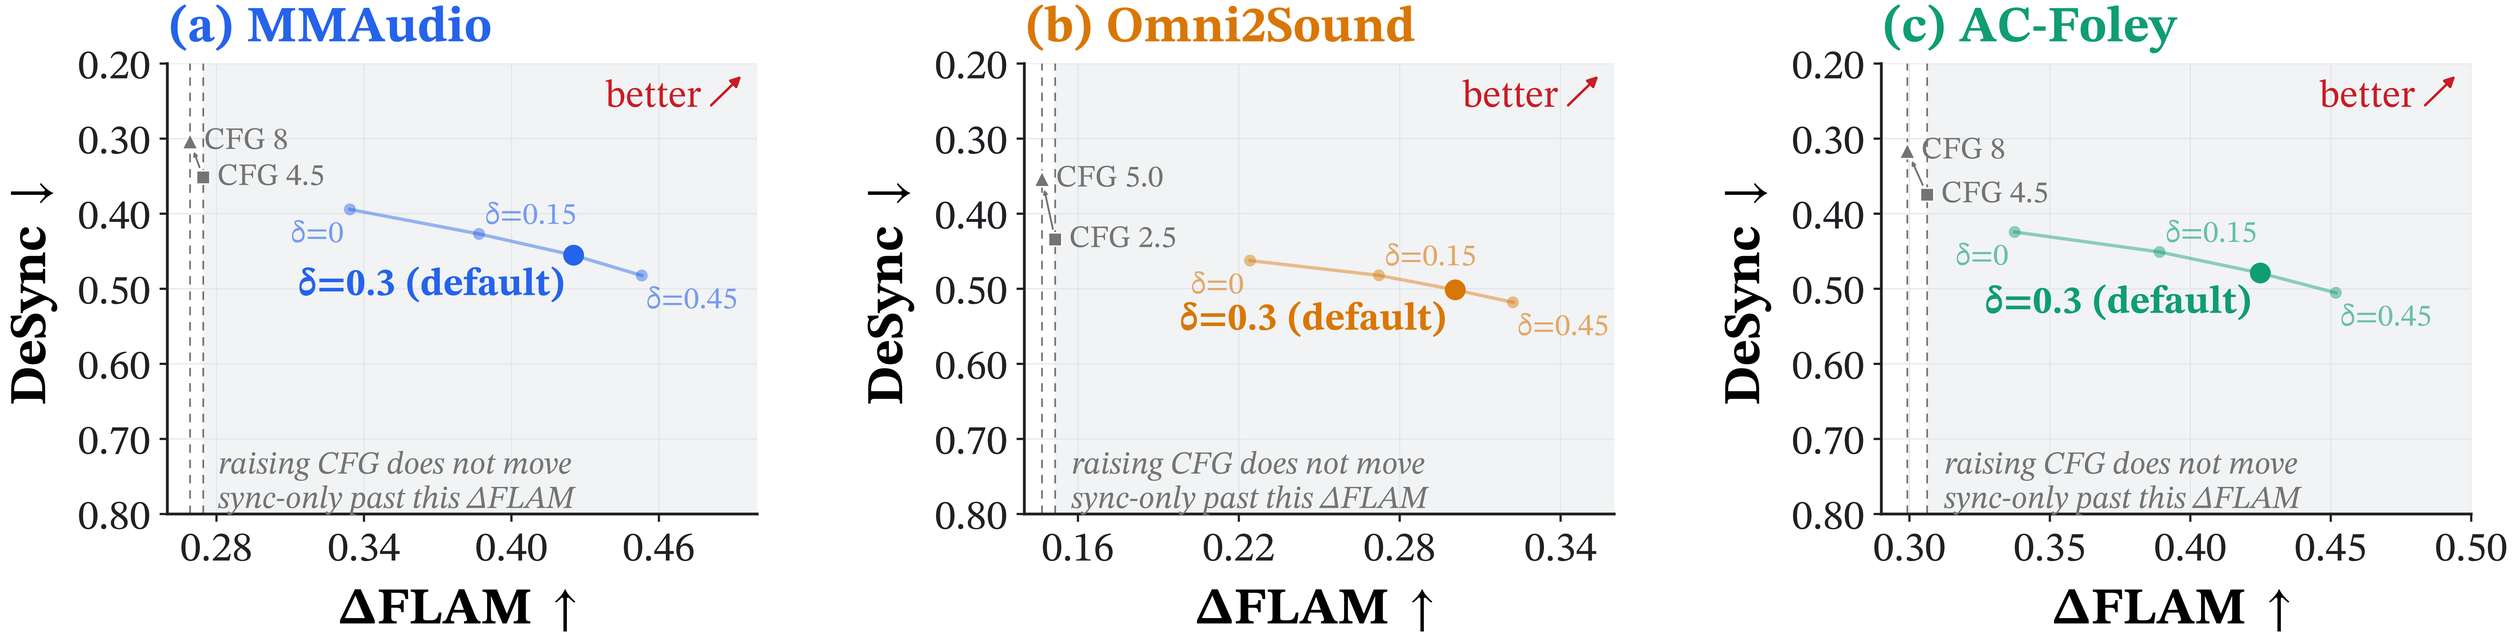

■ sync-only, default CFG ▲ sync-only, higher CFG ●—● TAP-Foley (δ = 0–0.45; default δ = 0.3)

Sync-only: semantic (CLIP) video features masked, synchronization features kept. Colored marks follow the backbone color. Shaded: ΔFLAM beyond the sync-only line.